

Supplementary material

How much does the map of science depend on the embedding model?

Alejandro Treny Ortega • Independent researcher

18 September 2026

This supplement gives the operational definitions, sensitivity checks and full result inventory for the manuscript. Each numbered table comprises a typeset summary and complete, unrounded CSV files in its matching `tables/data/Sxx/` directory. Figure data are supplied separately under `figures/data/`. The summaries do not replace those files.

Repeated selections come from the same frozen corpus. Their ranges describe sensitivity to those selections, not population confidence intervals. Model agreement does not establish thematic correctness. The chronology below distinguishes initial choices from later exploratory checks; the retained alerts and source problems are part of the results.

Contents

S1 Corpus, models and provenance	2
S2 Agreement measures and sampling stability	6
S3 Text input, pooling and truncation	9
S4 Scale, neighbourhoods and concrete cases	11
S5 Time, discipline composition and model traits	16
S6 Geometric properties and construct checks	21
S7 All Field-pair comparisons	26
S8 Reproducibility and access	31

S1 Corpus, models and provenance

S1.1 Eligibility, quotas and quality flags

OpenAlex’s core catalogue was queried for 2000–2024 records of type article, review or conference-paper, with an English language label, an abstract, no retraction or paratext flag and a known primary Field. Local preparation required a non-empty title and at least 50 regex words in the reconstructed abstract. It preserved literal content rather than translating or generating missing summaries. The unit was a distinct OpenAlex work ID, not a deduplicated intellectual work.

Seeded random request blocks supplied the general cohort. Saved candidates were processed in their recorded order until 400,000 eligible, distinct IDs were retained. The 100,000-record supplement was assigned to the smallest of the 130 Field-period groups, with deterministic Field/year tie-breaking. This raised minimum coverage without reducing the larger groups. Exact year counts were retained rather than forcing equal quotas within each five-year period. Raw retrieval ended on 15 September 2026; API responses were collected over time, so the corpus is not a simultaneous database snapshot. Saved responses, ID order and content hashes are required to reconstruct it; seeds alone would not reproduce a changing API.

Cleaning rejected recognised non-abstract content, explicit withdrawal/retraction titles and clear language failures. For language, fastText’s lid.176 model was applied to the whole abstract and three consecutive parts. A record was excluded only if the whole text had a non-English top label with score at least 0.9 and all three parts agreed on that language with scores at least 0.8. These are conservative operational scores, not calibrated probabilities or independent tests. Remaining uncertainty, possible mixed language, 50–79-word or unusually long abstracts, possible notices and duplicate DOI/text were marked. Content rules were developed by inspecting the downloaded corpus before embedding; they were not preregistered before extraction.

The final corpus retains 9,883 abstracts with language uncertainty, including 3,002 flagged as potentially mixed. These sets overlap. It also includes 42,849 abstracts of 50–79 words and 249 over 2,000 words. Subsequent analysis uses the flags for sensitivity rather than silently changing the corpus. Source checks establish agreement with downloaded records, not that every record is correct. Table S1 supplies the distribution and the full flag tables.

Table S1: Clean corpus by Field, period and cohort.

Field	2000-04	2005-09	2010-14	2015-19	2020-24	Base	Extra
Agricultural and Biological Sciences	2,166	2,528	3,940	4,298	6,491	18,920	503
Arts and Humanities	2,166	2,166	2,186	2,732	3,380	10,696	1,934
Biochemistry, Genetics and Molecular Biology	2,739	4,071	5,251	5,596	6,840	24,497	0
Business, Management and Accounting	2,166	2,166	2,186	2,660	4,040	11,137	2,081
Chemical Engineering	2,166	2,166	2,166	2,166	2,166	1,195	9,635
Chemistry	2,166	2,166	2,166	2,166	2,166	7,290	3,540
Computer Science	3,310	5,594	7,094	6,858	9,972	32,828	0
Decision Sciences	2,166	2,166	2,166	2,166	2,166	3,817	7,013
Earth and Planetary Sciences	2,166	2,166	2,166	2,166	2,166	7,256	3,574
Economics, Econometrics and Finance	2,166	2,166	2,166	2,166	2,845	8,620	2,889
Energy	2,166	2,166	2,166	2,166	2,166	2,617	8,213
Engineering	6,990	11,310	15,203	14,738	18,907	67,148	0
Environmental Science	2,166	2,876	4,005	4,018	5,919	18,569	415
Immunology and Microbiology	2,166	2,166	2,166	2,166	2,166	4,456	6,374
Materials Science	2,166	2,166	2,962	3,152	3,787	13,335	898
Mathematics	2,166	2,166	2,166	2,166	2,166	4,581	6,249
Medicine	7,128	10,946	15,771	18,629	26,195	78,669	0
Neuroscience	2,166	2,166	2,166	2,166	2,166	6,824	4,006
Nursing	2,166	2,166	2,166	2,166	2,166	2,077	8,753
Pharmacology, Toxicology and Pharmaceutics	2,166	2,166	2,166	2,166	2,166	1,715	9,115
Physics and Astronomy	2,166	2,704	3,094	3,388	3,470	14,772	50
Psychology	2,166	2,166	2,166	2,577	3,970	10,816	2,229
Social Sciences	2,776	4,893	7,110	8,891	13,879	37,549	0
Veterinary	2,166	2,166	2,165	2,165	2,165	642	10,185
Dentistry	2,165	2,165	2,165	2,165	2,165	2,238	8,587
Health Professions	2,165	2,165	2,165	2,165	2,833	7,736	3,757
Total	68,427	83,908	101,289	107,858	138,518	400,000	100,000

Total: 500,000 distinct OpenAlex work IDs. Period columns contain base plus extra; the last two columns divide those same articles and must not be added again. Core catalogue; years 2000–2024; article, review and conference-paper; OpenAlex English label, non-empty title, reconstructible abstract with at least 50 regex words, known primary Field, no retraction or paratext flag. Additional cleaning rejects clear language/content failures and marks ambiguous cases. No claim that every retained abstract is exclusively English or correctly labelled. Exact annual/cohort counts and quality flags are included as CSVs; raw retrieval ended on 15 September 2026. The database was not a simultaneous snapshot.

S1.2 Checkpoints, input and truncation

All ten models use the same ordered IDs and original title/abstract text. The customary separators, full model revisions and any adapter revisions are listed in Table S2. Outputs were saved as 32-bit vectors; subsequent geometric calculations use 64-bit arithmetic. Missing or degenerate outputs were checked, and each run records code, configuration, file and environment hashes.

Mean pooling averages every non-padding last-layer position, including model-added special tokens. CLS is the first-position state before any pooler; SEP is the final non-padding position, rather than the separator between title and abstract. Only BERT, SciBERT, BioBERT and PubMedBERT change under the CLS/SEP scenarios. SimCSE uses its pre-pooler CLS state; SPECTER2 uses its proximity adapter. MPNet and MiniLM retain their sentence-transformer mean-pooling configuration. Readout scenarios are not additional independently trained models.

The common-fragment comparison finds a shared prefix of content characters that fits every model, keeping each tokenizer and its customary separators. It uses a separate 52,000-ID panel and compares that fragment with the native representation of those same IDs. It does not equate token sequences between vocabularies. MiniLM’s separate 512-token check uses its supported architectural limit; vectors for texts unaffected by the original truncation are reused exactly. Neither check demonstrates that longer input gives better thematic representations.

Table S2: Frozen checkpoints, text formats and native truncation.

Model	Checkpoint and full revision	Input / truncation
SPECTER	allenai/specter	Title [SEP] abstract
	Revision: 81cbfb43d4fc2e728d5b4201ce14987db8d	Limit: 512
	0854c	Truncated: 5.66%
SPECTER2	allenai/specter2_base	Title [SEP] abstract
	Revision: 3447645e1def9117997203454fa4495937b	Limit: 512
	fbd83	Truncated: 5.00%
	Adapter: allenai/specter2	
SciNCL	2081559630a80fc5851d8f798a05ba81e9468089	
	malteos/scincl	Title [SEP] abstract
	Revision: ebc5348d184ba2fc9beee69b4e394263fce	Limit: 512
SciBERT	57b2e	Truncated: 5.00%
	allenai/scibert_scivocab_uncased	Title, blank line,
	Revision: 24f92d32b1bfb0bc9ab193ff3ad01e877	abstract
BERT	32fc1	Limit: 512
		Truncated: 4.96%
	google-bert/bert-base-uncased	Title, blank line,
	Revision: 86b5e0934494bd15c9632b12f734a8a67f7	abstract
MPNet	23594	Limit: 512
		Truncated: 6.87%
	sentence-transformers/all-mpnet-base-v2	Title, blank line,
MiniLM	Revision: e8c3b32edf5434bc2275fc9bab85f82640a	abstract
	19130	Limit: 384
		Truncated: 20.16%
	sentence-transformers/all-MiniLM-L6-v2	Title, blank line,
PubMedBERT	Revision: 1110a243fdf4706b3f48f1d95db1a4f5529	abstract
	b4d41	Limit: 256
		Truncated: 49.58%
	microsoft/BiomedNLP-BiomedBERT-base-uncased	Title, blank line,
BioBERT	-abstract-fulltext	abstract
	Revision: e1354b7a3a09615f6aba48dfad4b7a613ee	Limit: 512
	f7062	Truncated: 4.75%
SimCSE	dmis-lab/biobert-v1.1	Title, blank line,
	Revision: 551ca18efd7f052c8dfa0b01c94c2a8e68b	abstract
	c5488	Limit: 512
SimCSE		Truncated: 9.00%
	princeton-nlp/unsup-simcse-bert-base-uncas	Title, blank line,
	ed	abstract
SimCSE	Revision: 218baf6bc48c463a53f5b4e1c568299537f	Limit: 512
	ab124	Truncated: 6.87%

Truncation is measured on the full 500,000-article title-plus-abstract corpus and is not an error rate. The principal MiniLM limit remains 256 tokens; its 512-token variant is a separate sensitivity. CLS/SEP alternatives apply only to BERT, SciBERT, BioBERT and PubMedBERT. SEP takes the last non-padding position, not the title separator. Tokenizer differences mean that equal token limits would not imply equal text. Model-card URLs and all revisions are in the CSV.

S2 Agreement measures and sampling stability

S2.1 Coverage and chronology

Table S3: Coverage and chronology of the comparison blocks.

Block	Coverage	When specified	Interpretation
Initial comparison	500,000 corpus; 130 cells; 45 fixed model pairs	Before this phase's agreements	Normalised vectors; corrected CKA, k25; CLS/SEP and size checks.
Common fragment	52,000 paired IDs, distinct from the input panel	Before first scientific comparison	Same content characters; native tokenizers and separators.
Centres / random groups	256 per centre; 26/217/26 centres	Added after initial results	Matched dates/sizes; addresses centre averaging, not thematic validity.
Input expansion	26,000 to 52,000 nested IDs	After pilot; before expanded result	Three inputs and pooling crossed; 100 comparable selections.
Matched local searches	217 groups, 50 queries, ten selections	Expanded phase design	256 candidates per scope, including identical queries and matched dates.
Atlas	217 Subfields; 183 in sensitivity panel; 12 paper examples	After global results; before reading selected titles	Rule-selected examples; source errors retained.
Morphology	52,000 IDs; alternative measures and synthetic clouds	Exploratory extension after core results	Protocol before new real-data metrics. Neighbour-fraction check added during execution.
Field pairs	325 pairs per property; 26 saved conditions	After morphology results	Direction/magnitude rules fixed before new pair counts; no new inference.

This is a sequence of documented phases, not a preregistered study. Sources: ANALYSIS_PROTOCOL, ROBUSTNESS_PROTOCOL, CASE_ATLAS_PROTOCOL, MORPHOL-

OGY_PROTOCOL and FIELD_PAIR_PROTOCOL, with their dated decision entries. Later controls are not relabelled as initial hypotheses.

S2.2 Alignment and neighbour definitions

For each comparison, rows of X and Y identify the same articles. Vectors are scaled to unit length, and columns are then centred for CKA and Procrustes. Let n be the row count and define

$$\begin{aligned} S_{XY} &= \|X^\top Y\|_F^2, & R_{XY} &= \sum_i \|X_i\|^2 \|Y_i\|^2, \\ T_X &= \sum_i \|X_i\|^2, & H_{XY} &= S_{XY} - \frac{n}{n-2} R_{XY} + \frac{T_X T_Y}{(n-1)(n-2)}. \end{aligned}$$

The reported corrected alignment is

$$\text{CKA}_{\text{corr}}(X, Y) = \frac{H_{XY}}{\sqrt{H_{XX} H_{YY}}}.$$

The common HSIC factor cancels. Negative corrected values are retained, and degenerate inputs are rejected. The implementation was checked independently using relation matrices with zero diagonals. The correction addresses finite-sample inflation (Kornblith et al., 2019; Murphy et al., 2024); the complete ratio is not asserted to be exactly unbiased.

Procrustes similarity is $\|X^\top Y\|_*/\sqrt{T_X T_Y}$, where the numerator is the sum of singular values. Its arccosine gives the reported angular distance. Rotation/reflection and a global scale can be aligned; an arbitrary deformation cannot. Unequal embedding dimensions can be understood by adding zero coordinates to the smaller representation. Normalisation precedes centring, so this procedure does not imply invariance to translating the original vectors and normalising them again (Williams et al., 2021).

The distance-rank check correlates cosine distances for 20,000 distinct article pairs selected with a recorded seed, identically in both representations. Ties receive average ranks. These pair observations and the alternative alignment measures are not treated as independent replications (Williams, 2024).

For query i and models a, b , neighbour overlap is

$$O_k(i; a, b) = \frac{|N_k^a(i) \cap N_k^b(i)|}{k}.$$

The query is excluded by ID. In the core exact search, similarities are computed in 64-bit blocks and rounded to twelve decimal places for deterministic tie resolution by global row order. Independent sorting checks verify the stored lists. The combinatorial reference for independent uniform lists among M eligible candidates is k/M ; the adjusted overlap

is $(O_k - k/M)/(1 - k/M)$. This reference does not remove density, popular-neighbour or composition effects. The primary results therefore retain raw overlap alongside explicitly matched candidate pools.

S2.3 Selection screens and retained alerts

The original screen uses twenty selections at each of 1,024 and 2,048 articles per Field-period group, paired between models. It flags an absolute median shift above 0.02 or a central 95% selection width above 0.04 at the larger size. A row permutation per model/group provides an alignment diagnostic rather than a permutation-test p -value. The quality sensitivity retains one representative per connected duplicate DOI/text component and excludes specified flags, leaving 448,886 records. It changes the operational population and does not replace the main corpus.

Subfield screens use twenty nested selections up to 256/512 articles, subject to availability. Groups below eight articles lack a two-size diagnosis. If the larger draw equals the full group, its zero selection width is explicitly marked and cannot establish precision for a wider population. Table S4 keeps these cases separate from the original Field-period alerts.

Table S4: Agreement of alternative measures and retained stability alerts.

Measures compared	Cells	Mean rank agreement	Minimum	Maximum
Corrected CKA / Procrustes	130	0.891	0.793	0.960
Corrected CKA / Distance ranks	130	0.881	0.777	0.963
Procrustes / Distance ranks	130	0.789	0.649	0.893

The top panel reports descriptive Spearman correlations of model-pair agreements within each of 130 Field-period cells. These correlations compare measures; they do not validate thematic content. All 41 saved summaries and all individual original alerts accompany the table.

- Original Field-period screen: 50 alerts among 5,850 comparisons; largest central-95% selection width among flagged cases: 0.064.
- The larger sample does not equal the entire group: 2,628 alerts / 8,235 comparisons; 183 Subfields.
- The larger sample equals the entire group: 213 alerts / 3,060 comparisons; 68 Subfields.

A zero selection width when the whole group is used does not establish stability at a smaller size. These are within-corpus screens, not population confidence intervals. Alerts from different designs remain separate.

S3 Text input, pooling and truncation

S3.1 Model-specific input effects

The expanded panel uses 400 articles per Field-period group, 2,000 per Field, selected by a fixed ID-based order. It retains the 26,000 pilot IDs. For each model, the model-change effect is one minus mean agreement with the other nine models on title plus abstract. The input-change effect is one minus agreement between that model’s combined input and its title-only or abstract-only input. Their difference is positive when model change is larger. Neighbour changes use the same 2,000 candidates per Field; self is excluded.

A symmetric summary of input transitions checks dependence on the combined-input reference. Mean, CLS and SEP scenarios cross all three inputs; only the four word-level BERTs change readout, which also changes the cross-model reference. Saved combined-input and pilot vectors were reused with exact ID/text alignment and bitwise checks. The analysis makes a practical comparison of input conditions, not a causal decomposition of architecture and content.

Table S5: Input changes by model, averaged equally over the 26 Fields.

Model	Title: shape	Abstract: shape	Title: neighbours	Abstract: neighbours
SPECTER	0.208	0.044	61.1%	31.3%
SPECTER2	0.203	0.048	60.6%	31.3%
SciNCL	0.205	0.028	57.3%	23.9%
SciBERT	0.532	0.014	84.2%	15.8%
BERT	0.418	0.013	83.6%	18.0%
MPNet	0.189	0.022	51.7%	17.7%
MiniLM	0.246	0.031	58.1%	21.5%
PubMedBERT	0.526	0.017	82.9%	16.2%
BioBERT	0.636	0.024	83.4%	19.2%
SimCSE	0.496	0.028	83.8%	26.3%

Same 52,000 IDs and 2,000 candidates per Field. Shape change: 1 minus corrected CKA. Neighbour change: fraction of the 25 neighbours replaced. Title/abstract conditions compare each model with its own full input. Model-choice means and all Field-specific values are included as CSVs, alongside the full recipe cross. These are practical input changes with native text limits, not a causal separation of encoder and text length.

Title-only model-minus-input contrast (positive: changing model has the larger mean effect): MEAN, shape = 0.0042; MEAN, neighbours = -0.0191; CLS, shape = 0.0363; CLS, neighbours = 0.0304; SEP, shape = 0.0627; SEP, neighbours = 0.0171. Only the four word BERTs change readout in CLS/SEP. Close means do not establish equivalence.

S3.2 The 26k to 52k stability check

Both sizes were assessed with 20 and then 100 selections without replacement, using 500 articles per Field in the pilot and 1,000 in the expansion, balanced across dates. The first twenty repeats reproduce the earlier diagnostic. The same 0.02 median-shift and 0.04 central-width rules apply to the principal shape contrast. Retaining the original 31 keys distinguishes persistent, resolved and newly appearing alerts. The 100-selection comparison yields 72 to five alerts, while the historical twenty-selection comparison yields 31 to four. These counts must not be substituted for one another. They do not extend automatically to every readout or neighbour analysis.

Table S6: Input-alert transitions for the nested 26k and 52k panels.

Selections / scope	Contrasts	26k alerts	52k alerts	Persist	Resolved	New
20 / model-Field	520	31	4	3	28	1
20 / Field average	52	0	0	0	0	0
100 / model-Field	520	72	5	5	67	0
100 / Field average	52	0	0	0	0	0

The five remaining model-Field-input alerts with 100 selections are: BioBERT / Agricultural and Biological Sciences / abstract (width 0.047); BioBERT / Earth and Planetary Sciences / abstract (width 0.045); BioBERT / Economics, Econometrics and Finance / abstract (width 0.042); BioBERT / Immunology and Microbiology / abstract (width 0.063); PubMedBERT / Health Professions / title (width 0.040). No alert does not imply an effect different from zero. Original 20-selection and later 100-selection counts are kept separate.




Figure S1: Input stability alerts for nested samples. Counts are model–Field–input contrasts out of 520. The original 20-selection comparison and later 100-selection check use the same recorded thresholds, but their counts are not interchangeable. These are selection-sensitivity checks within the fixed corpus. The 52 Field-averaged contrasts are reported separately in Table S6.

S4 Scale, neighbourhoods and concrete cases

S4.1 Matched scopes and candidate sizes

Within-group comparisons use 128, 256 and 512 articles, with nested ID-based selections and the same 183 eligible Subfields at all three sizes. The native analysis retains all 252 Subfields where the measure is defined; groups with exactly $k + 1$ articles have uninformative complete neighbour lists and are marked accordingly. A separate time comparison retains the same 125 Subfields with at least 128 articles in every period.

Centre comparisons use means of unit article vectors. Direct and weighted-Subfield reconstructions of Field centres were checked for equality. Matched centres contain 256 articles per group; random allocations preserve counts by period and are shared across models. Twenty selections of one Subfield per Field restrict the centre count to 26. The random reference retains cross-model article correspondence, so a positive agreement is expected and is not a null test of independence.

The paired search fixes 50 queries per eligible Subfield, with 256 candidates in each scope and ten selections. The same 50 queries are included in both candidate sets; only the other 206 positions change scope. Date counts match. One query per model, selection, group and

scope was independently sorted, giving 43,400 checks. Results for those 10,850 queries have repeated candidate controls; local scores for all 500,000 records do not.

Table S7: Matched Field/Subfield neighbour searches.

k	Field overlap	Subfield overlap	Difference (pp)	Lower / higher	Subfields
10	38.7%	36.9%	-1.73	145 / 72	217
25	46.7%	45.3%	-1.35	127 / 90	217
50	54.5%	53.7%	-0.79	109 / 108	217

Difference is Subfield minus Field. Each scope contains 256 candidates, including the same 50 query articles, with matched dates; self is excluded during search. Means use ten candidate selections. The broader Field comparison is conditioned on the queries and is not 256 unrestricted Field draws. All 217 Subfields appear in the paired design; 183 meet all size conditions of the separate 128/256/512 sensitivity panel. Full size/readout/k results accompany the table.

S4.2 Subfield persistence

The atlas compares ranks among the same 183 Subfields under 27 combinations of three candidate sizes, three readouts and three neighbour counts. Example rules were fixed after the overall results but before reading selected article titles. The six upper and six lower examples are descriptive extremes under those rules. Their rank ranges locate persistent relative agreement, not semantically correct or incorrect Subfields.




Figure S2: Relative neighbourhood-agreement ranks for twelve rule-selected Subfield examples. Ranges span 27 combinations of candidate size, readout and neighbour count, using the same 183 eligible Subfields. The upper/lower six illustrate the recorded higher/lower agreement rule. These are sensitivity ranges, not confidence intervals; higher agreement does not imply better thematic accuracy.

S4.3 Article and relation examples

The article atlas uses ten candidate selections for the fixed queries. Link comparisons are restricted to 531,650 directed relations whose endpoints remain eligible, avoiding the error of treating an unavailable target as a rejected neighbour. For each displayed source, both persistent and model-dependent targets are retained in the accompanying files. Figure S3 selects the eligible relation with maximal model disagreement for that source; it cannot estimate how common disagreement is.

The examples were inspected against saved OpenAlex records. The history-book, infection/politics and announcement cases remain visible, as does one target with an incomplete abstract. A diagnostic for 82 generic titles removes five fixed queries and barely changes

the overall query mean (45.344% to 45.349%). That narrow check neither estimates the prevalence of thematic classification errors nor cleans all candidates.

Table S8: All twelve rule-selected article examples, including source caveats.

Article and OpenAlex ID	Subfield label	Mean	Selection range	Inspection note
SIRT6 overexpression induces massive apoptosis in cancer cells but not in normal cells W2022976570; 2011	Geriatrics and Gerontology	79.3%	76.7%–82.4%	No clear issue in this limited inspection; not expert validation.
Comparison of The Effectiveness of Scientific Approach and Problem-Solving Approach in Problem-Based Learning in Class IX of SMP Negeri 3 Pangsid W2998141762; 2019	Applied Mathematics	84.1%	75.1%–87.4%	Mathematics education under Applied Mathematics; administrative label.
A systematic analysis on the managerial non-efficiency of the state-owned businesses W2359252036; 2002	General Decision Sciences	76.9%	74.7%–78.8%	No clear issue in this limited inspection; not expert validation.
Application of multi-agent approach to development of the petrol dispatching system of gas stations network W3107610255; 2020	Discrete Mathematics and Combinatorics	78.2%	74.0%–82.4%	No clear issue in this limited inspection; not expert validation.
Persuasive Talk, Delegation, and Announcements W2252942328; 2012	Safety Research	78.2%	72.1%–83.3%	Persuasion economics under Safety Research; label caveat.
Road risk-perception and pedestrian injuries among students at Ain Shams University, Cairo, Egypt W2081778888; 2011	Safety, Risk, Reliability and Quality	76.4%	70.8%–85.2%	No clear issue in this limited inspection; not expert validation.

Continued on next page

Table S8 (continued)

Article and OpenAlex ID	Subfield label	Mean	Selection range	Inspection note
Invasion Strategy W2011180849; 2003	Political Science and International Relations	18.7%	16.1%– 21.0%	Bacterial-infection abstract labelled as politics.
Beyond the witch trials: Witchcraft and magic in Enlightenment Europe W7045581866; 2004	Nuclear and High Energy Physics	21.4%	19.2%– 23.7%	History-book description labelled as physics; type caveat.
Optical properties of cobalt clusters implanted in thin silica layers W1966552042; 2006	Ceramics and Composites	23.2%	21.3%– 24.4%	Substantial formula code in abstract; input caveat.
Announcement W4234898602; 2010	Ophthalmology	22.2%	19.9%– 25.8%	Announcement / atlas review; not detected by the original notice screen.
Conversion of a rule based to an object oriented expert system W2117277818; 2002	Software	25.0%	22.8%– 26.0%	No clear issue in this limited inspection; not expert validation.
‘Thare is na leid on life of lordship hym like’: linguistic means of depicting king arthur in the poem Golagros and Gawain W3110667008; 2020	Classics	25.5%	24.4%– 26.3%	No clear issue in this limited inspection; not expert validation.

Numbers are k25 overlap averaged over 45 model pairs; ranges span ten candidate selections. The first/last six examples come from high/low agreement sources under the recorded rule. Titles were read after selection. Full source/target IDs, titles and per-model supports for both relation types are in the accompanying CSVs. The most model-dependent link was selected for each source; this does not estimate prevalence. An incomplete target abstract is also retained and marked in Figure S3. None of these checks validates OpenAlex labels.

SIRT6 / cancer [b]	10	9	10	10	10	10	10	0	1	10
Mathematics education [a]	4	7	4	0	8	5	7	0	7	10
State-owned businesses	9	2	3	10	10	8	0	10	10	10
Petrol dispatch	10	1	0	10	10	2	7	10	10	10
Persuasive talk [a]	4	10	10	10	4	8	4	10	0	10
Pedestrian injuries	10	8	3	10	10	1	0	10	10	9
Invasion Strategy [a]	0	10	10	0	10	10	10	0	3	0
Witch trials [a]	0	10	0	10	10	10	0	10	10	10
Cobalt / silica layers	10	10	10	0	0	10	10	0	0	10
Announcement [a]	0	0	8	10	10	0	0	10	0	10
Expert-system conversion	0	0	0	10	10	10	10	10	10	1
Arthurian poetry	0	0	10	10	10	9	10	10	10	10
	SPECTER	SPECTER2	SciNCL	SciBERT	BERT	MPNet	MiniLM	PubMedBERT	BioBERT	SimCSE

Figure S3: Most encoder-dependent eligible link for each of the twelve fixed article sources. Each cell counts selections in which the target is among the 25 nearest neighbours (0–10); both endpoints are always eligible within the 256-candidate design. Upper/lower six rows are high/low agreement sources. [a] Source or label caveat; [b] the selected target has an incomplete abstract. Full IDs, titles and caveats are retained in Table S8’s files. Links were selected for maximal model disagreement, so this panel does not estimate its prevalence.

S5 Time, discipline composition and model traits

S5.1 Period contrasts and search opportunities

Period summaries give equal weight to Fields and fixed model pairs. Endpoint differences compare 2020–24 with 2000–04; the full five-period sequence determines whether the trajectory is monotone. Native candidate pools and 2,048-candidate pools are reported separately. A later identical-query comparison separates changes in selected queries from candidate count. Publications are represented by the same present-day models across all dates; differences may reflect coverage, writing, composition, text length or model exposure rather than historical changes in scientific organisation.

Table S9: Temporal endpoint changes and Field heterogeneity.

Measure / candidate design	Mean change	Positive	Negative	Non-monotone
CKA, native groups	0.0207	23	3	22
CKA, 2,048 per group	0.0206	24	2	22
Distance ranks, native	0.0304	22	4	19
Procrustes, native	0.0019	15	11	24
k25, all candidates	-0.0087	10	16	22
k25, 2,048 / mean	0.0091	20	6	20
k25, 2,048 / CLS	0.0034	16	10	22
k25, 2,048 / SEP	0.0138	21	5	20

End minus start: 2020–24 versus 2000–04. Each row describes 26 Fields with equal weight; neighbour changes are fractions, so multiply by 100 for percentage points. Five periods are used to classify monotonicity. Full k10/k25/k50, recipe, identical-query and pair-level checks are in the CSVs. Candidate counts can change the sign. These are descriptive period contrasts, not causal historical convergence.




Figure S4: Descriptive time patterns under candidate controls. Values are changes from 2000–04. Fields and fixed model pairs have equal weight. In (A), thin lines show the 26 Fields; the thick line is their mean. In (B), k25 overlap is compared with all candidates or 2,048 candidates. Changes in search opportunity can reverse the average sign. These are not causal historical trajectories.

S5.2 Medicine and Subfield composition

Composition controls draw 2,048 articles per Field with identical date quotas under observed or balanced Subfield composition, repeated ten times. They are also summarised after excluding BioBERT and PubMedBERT. Energy has only one eligible Subfield in this control, so its composition does not change in the same way as Medicine’s. Balancing and model omission are descriptive sensitivities, not interventions isolating independent causes. The complete 26-Field results accompany the four displayed examples.

Table S10: Discipline agreement with observed and balanced Subfield composition.

Field	Models	CKA obs.	CKA bal.	k25 obs.	k25 bal.
Medicine	Ten	0.546	0.543	29.0%	29.2%
Medicine	Eight	0.564	0.562	29.5%	29.9%
Energy	Ten	0.772	0.772	33.2%	33.2%
Energy	Eight	0.774	0.774	33.4%	33.4%
Mathematics	Ten	0.671	0.644	29.4%	28.8%
Mathematics	Eight	0.691	0.675	30.8%	30.3%
Physics and Astronomy	Ten	0.663	0.674	33.0%	34.1%
Physics and Astronomy	Eight	0.665	0.671	33.3%	34.3%

Eight excludes BioBERT and PubMedBERT. 2,048 articles per Field with identical date quotas; mean of ten selections. Energy has one eligible Subfield, so balancing does not introduce a comparable composition change there. The display focuses on the four Fields of Figure S5; all 26 Fields and matched contrasts are supplied. Composition and biomedical training are not isolated causal explanations.




Figure S5: Observed and balanced Subfield composition in four Fields. Each condition uses 2,048 articles per Field and identical date quotas. Bars are means and whiskers the range of ten selections, not confidence intervals. The eight-model panel excludes BioBERT and PubMedBERT. *Energy has one eligible Subfield. Balancing composition and omitting biomedical models do not isolate a unique causal explanation. All 26 Fields are retained in Table S10's files.

S5.3 Readout and apparent model families

Recorded model traits cover broad training-domain and objective groups. Pairwise models, leave-one-model checks and readout scenarios assess their descriptive association with agreement. Shared model membership makes the 45 pairs dependent. Some traits are linked in this panel, and coefficients can become non-identifiable under omission; for example, the SEP scenario without SimCSE retains the missing coefficients rather than reporting spurious estimates. The heatmaps show how apparent families depend on the representation recipe.

Table S11: Descriptive model-trait fits and identifiability.

Outcome	Rank	Condition no.	R squared	Unidentified coefficients after omission
CKA / primary	7/7	7.91	0.741	0
CKA / 2,048 articles	7/7	7.91	0.741	0
CKA / quality filter	7/7	7.91	0.752	0
CKA / CLS	7/7	18.59	0.459	14
CKA / SEP	7/7	8.33	0.258	7
Neighbours k25 / mean	7/7	7.91	0.759	0
Neighbours k25 / CLS	7/7	18.59	0.507	14
Neighbours k25 / SEP	7/7	8.33	0.442	7

Ten fixed models yield 45 overlapping pairs. R squared describes these pairs, not predictive validity. The final column counts unavailable coefficient rows across the ten leave-one-model checks, not independent models or tests. In the SEP design, omitting SimCSE makes coefficients non-identifiable; those values remain missing. Architecture, lineage, final corpus, objective, domain and readout are correlated traits. Full trait definitions, source URLs, coefficients, pair-label permutation references and omissions accompany this table. No causal training effect is claimed.




Figure S6: Pairwise shape agreement under three readout choices. Matrices show mean corrected CKA over the original 130 Field-period cells. Only BERT, SciBERT, BioBERT and PubMedBERT change readout; diagonal cells are identities. Model order in columns: 1 SPECTER, 2 SPECTER2, 3 SciNCL, 4 SciBERT, 5 BERT, 6 MPNet, 7 MiniLM, 8 PubMedBERT, 9 BioBERT, 10 SimCSE. Similarity patterns are not an isolated effect of shared training.

S6 Geometric properties and construct checks

S6.1 Properties, alternatives and selection stability

The geometric analysis reuses the 52,000-article input panel, with 2,000 articles per Field and equal period counts. All vectors have unit length and calculations use 64-bit arithmetic. For distinct vectors u_i, u_j , angular spread is the median of $\arccos(u_i^\top u_j)$, reported in degrees. Alternatives are its 10th/90th percentiles and the median angle to the centre direction. They describe spread rather than semantic diversity or isotropy (Rudman et al., 2022).

After centring the unit-vector cloud, let λ_j be the covariance eigenvalues and $p_j = \lambda_j / \sum_l \lambda_l$. The primary effective dimension is $1 / \sum_j p_j^2$, equivalent to the participation ratio in the main text. The entropy alternative is $\exp(-\sum_j p_j \log p_j)$, applied to the covariance spectrum (Roy and Vetterli, 2007). D80 counts the smallest number of descending eigenvalues containing 80% of variance. These quantities are unchanged by adding zero coordinates, so the 384-versus 768-coordinate comparison is not normalised by ambient dimension. They describe linear variation, not nonlinear intrinsic dimension or a count of topics.

Twenty half samples contain 1,000 articles per Field, and five further equal-size selections contain 2,000, all drawn from the fixed corpus with matched IDs across models. Central 90% ranges and absolute median shifts are divided by the primary value. Operational thresholds are 5% for angular spread, 10% for PR and 20% for connection. With five repeats, quantiles are coarse; individual values are retained. Passing those selections does not erase the four PR half-sample alerts. The complete analysis contains 16,884 sets of measures, including time, size and representation controls; not every alternative was repeated under every selection.

Table S12: Selection stability of the geometric properties.

Property	Selection design	Cases	Alerts	Median width	Maximum width
Angular spread	20 half samples	260	0	0.79%	1.81%
Angular spread	5 equal-size selections	260	0	0.49%	1.85%
Effective linear dimension	20 half samples	260	4	4.06%	12.84%
Effective linear dimension	5 equal-size selections	260	0	2.60%	7.53%
Connection	20 half samples	260	258	18.62%	423.07%
Connection	5 equal-size selections	260	51	10.59%	289.69%

A case is one model-Field combination. Width is the central 90% range divided by the primary value; the screen also checks median shift. Thresholds: 5% angular spread, 10% PR, 20% connection. These are operational sensitivity thresholds, not journal standards or population precision. The four PR half-sample alerts remain visible even though the five equal-size selections pass. PR is effective linear dimension, not number of topics. Spectral alternatives share the same covariance spectrum and do not independently validate thematic meaning.

S6.2 Sample size




Figure S7: Sample-size sensitivity of the three saved geometric properties. Each line shows one model’s median relative change across the 26 Fields, compared with 2,000 articles per Field. Panels use different vertical scales. Connection uses the union-graph spectral gap with $k = 25$ fixed as n changes; the separate approximately fixed-fraction check is supplied in Table S13’s files. These curves do not certify population precision.

S6.3 Representation and content controls

Controls vary quality flags, title/abstract input, readout, common content, extreme points, global centring and the MiniLM token limit. The quality comparison replenishes each period to its specified size after applying recorded flags. The extreme-point control removes the ten points farthest from the centre in each period of the initial half sample; removing ten random points per period supplies its equal-size reference. It does not presume that extreme points are erroneous.

Global centring subtracts the mean of the balanced 52,000 unit vectors and then normalises each result again. This changes the representation and is distinct from centring a covariance matrix for PR. Three Gaussian clouds per model–Field case provide geometric references with matching mean/covariance only in expectation before unit normalisation. Their realised moments are checked; these references are neither exact matched-moment controls nor calibrated statistical null distributions. The point controls in Figure S8 and Table S16 have their own matched references and must not be treated as another complete set of repeated selections.




Figure S8: Agreement of Field rankings before and after each matched control. Each point is a model’s Spearman rank agreement; horizontal spans show their range. CLS/SEP change only the four word BERTs. The dashed line at 0.9 is a visual reference, not a pass/fail threshold. Global centring changes the representation. All saved controls, including additional point sensitivities, accompany Table S12.

S6.4 Why connection is not general fragmentation

Connection uses the second-smallest eigenvalue of the normalised graph Laplacian $L = I - D^{-1/2}AD^{-1/2}$ on the unweighted union of directed nearest-neighbour links. The value is zero for a disconnected graph. Primary $k = 25$ is checked at 10 and 50, with a weighted local-scale graph and mutual-neighbour components as alternatives. Spectral-clustering theory motivates the construction (von Luxburg, 2007); it does not establish that a small gap means separate scientific themes. Eigenvalue residuals and independent dense calculations were checked.

Single-linkage merging supplies the radii at which the largest component reaches specified fractions of the cloud. Radius ratios and the maximum-to-median minimum-spanning-tree edge ratio provide different connection descriptions. Nine synthetic clouds, followed by Gaussian references and article controls, exposed non-equivalent responses to elongation, isolated outliers and separated components. Fixed k also changes its fraction of candidates as sample size grows. The approximately fixed-fraction comparison was added during execution and is identified as a later sensitivity. These failures are why connection is excluded from the manuscript’s interpretation of Field fragmentation.

Table S13: Construct checks on all nine saved synthetic clouds.

Cloud	Gap k25	Radius 90/50	Max/median edge	Mutual giant
One round cloud	0.499	1.121	1.371	96.6%
One narrow cloud	0.503	1.132	1.400	95.5%
One wide cloud	0.487	1.084	1.262	99.8%
One elongated cloud	0.091	1.115	1.299	99.1%
One low-rank cloud	0.098	1.589	4.913	99.9%
Two separated groups	0.000	1.766	2.253	49.5%
Four separated groups	0.000	1.003	1.587	24.9%
Two groups with bridges	0.007	1.005	1.428	97.8%
One cloud with outliers	0.499	1.123	4.770	95.7%

Each saved cloud has 1,000 points. Smaller gap can arise without separate groups; radius/edge ratios can respond to outliers or elongated clouds. In all 260 primary model-Field cases the union k25 graph is connected. Connection therefore remains a graph description and is not validated as general thematic fragmentation. The five equal-size selections retain 51 connection alerts (Table S12).

The 780 Gaussian reference rows are included. Observed-relative reference errors span -4.69% to 4.60% for angular spread, and -10.47% to 19.78% for PR. Gaussian draws preserve fitted moments only in expectation before renormalisation; they are not calibrated null tests. Full n/k controls accompany the table; the fixed-neighbour-fraction check was added during execution.




Figure S9: All nine saved synthetic-cloud checks. (A) Spectral gap of the union k25 graph; smaller means weaker connection. (B) Ratio of radii at which the largest component contains 90% and 50% of points. (C) Maximum divided by median minimum-spanning-tree edge length. Orange marks the elongated, low-rank and outlier clouds: apparent weak connection or late merging need not mean separate thematic groups. These are controlled geometric examples, not a calibrated statistical null.

S7 All Field-pair comparisons

S7.1 Directions and magnitude

All unordered Field pairs are enumerated by ascending OpenAlex ID. For each model and property, the 26 saved conditions are the primary selection, twenty half samples and five additional equal-size selections, with repeat indices matched across the two Fields. If $d_r = 2(B_r - A_r)/(A_r + B_r)$, direction is positive when $\min_r d_r > t + 10^{-10}$ and negative when $\max_r d_r < -t - 10^{-10}$. Otherwise it is unresolved. The minima $t \in \{0, 0.01, 0.05, 0.10\}$ were specified before the pair counts, after the morphology results. They are magnitude sensitivities, not tests of significance.

All-ten agreement requires every model to have the same persistent direction. Opposition requires at least one persistently positive and one persistently negative model. All remaining pairs are unresolved, including partial agreement. These classes are exhaustive and disjoint. The full range, rather than a favourable central interval, determines persistence; the overlapping selections do not constitute independent replications.

Table S14: All Field-pair conclusions under four minimum relative differences.

Property	Minimum	All ten agree	Opposition	Unresolved	Pairs
Angular spread	0%	42	262	21	325
Angular spread	1%	30	238	57	325
Angular spread	5%	10	96	219	325
Angular spread	10%	0	19	306	325
Linear dimension	0%	54	221	50	325
Linear dimension	1%	50	213	62	325
Linear dimension	5%	36	177	112	325
Linear dimension	10%	26	126	173	325

Every property uses all 325 Field pairs and ten models. A model direction must persist in all 26 saved conditions: primary sample, twenty halves and five equal-size selections. Opposition needs at least two fixed models with persistent opposite directions; unanimity needs all ten. Other cases are unresolved, not equal. The signed symmetric difference is $2(B - A)/(A + B)$; the minimum is an operational magnitude cutoff, not statistical significance or a share of science correctly represented. Full 650 pair-property rows and all 6,500 model directions accompany the table.




Figure S10: All 650 Field-pair/property classifications at zero minimum relative difference, using the 26 saved selection conditions. The upper triangle shows angular spread; the lower shows effective linear dimension (PR). Row labels and column numbers identify the same 26 Fields; names are shortened for display. Each triangle contains 325 pairs. Blue requires all ten models to agree persistently; orange requires at least two persistently opposing models; grey is unresolved. The diagonal is not a comparison. Full labels, model directions and magnitude checks are in Table S14's files.

S7.2 Alternative measures and model panels

An opposing model pair is retained under the alternatives only if those same two models preserve their original directions for every available alternative. Angular alternatives have 26 conditions; spectral alternatives have three (primary, first half sample and first additional selection). D80 is discrete and ties are unresolved. At positive minima, the principal property must exceed the minimum, while the alternatives must retain direction. Models are never replaced with different opposing models simply to preserve a count. Ten leave-one-model panels and the six similarity-oriented models test panel dependence; the smaller panel naturally offers fewer opportunities for opposition.

Table S15: Persistent opposition and identical witnesses across alternatives and model panels.

Model panel	Models	Spread	Spread + alternatives	PR	PR + alternatives
All ten	10	262	225	221	196
Six similarity models	6	231	190	160	121
Omit SPECTER	9	234	182	212	181
Omit SPECTER2	9	251	205	211	182
Omit SciNCL	9	259	217	215	190
Omit SciBERT	9	258	222	218	194
Omit BERT	9	261	223	211	183
Omit MPNet	9	262	223	218	192
Omit MiniLM	9	259	223	221	196
Omit PubMedBERT	9	261	220	206	188
Omit BioBERT	9	259	217	212	181
Omit SimCSE	9	259	217	217	191

All counts are out of 325 at zero minimum magnitude. Alternatives must retain the directions of the same opposing models; different witnesses cannot be substituted between measures. Angular alternatives have 26 saved conditions, spectral alternatives three. Six similarity models: SPECTER, SPECTER2, SciNCL, MPNet, MiniLM and SimCSE. Fewer models provide fewer opportunities for opposition; this panel does not isolate a training effect. All four magnitude cutoffs and witness identities are supplied as CSVs.

S7.3 Matched point controls

Each of the eleven point controls is compared with its appropriate equal-size reference. Changing quality, extremes or sample size changes IDs by design; other matched controls verify identical IDs. The summary distinguishes a changed class from survival of the same original opposing models and directions in both reference and variant. One favourable controlled comparison does not demonstrate persistence across 26 new selections. The saved differences were independently audited through equivalent ratio inequalities, covering 442,000 contrasts without importing the classification function.

Table S16: Point controls: retention of the original opposing models.

Control	Matched reference	Models changed	Spread retained	PR retained
Title only	half	10	246/262	193/221
Abstract only	half	10	262/262	221/221
Word BERTs: CLS	half	4	253/262	210/221
Word BERTs: SEP	half	4	257/262	201/221
Common fragment	common_native	10	262/262	216/221
Global centring	half	10	145/262	218/221
Quality flags	half	10	262/262	221/221
Extreme-point removal	trim_random	10	262/262	221/221
MiniLM 512	primary	1	262/262	221/221
n = 500	primary	10	262/262	219/221
n = 4,000	primary	10	262/262	221/221

Numerators retain the same original model pair with its original signs in both matched reference and variant. Denominators are the original persistent oppositions (262 angular spread; 221 PR). Values are not new 26-condition stability tests. The common fragment uses its own matched panel; extreme-point removal compares with random removal at equal size. Quality/extreme/size controls change IDs as specified; other matched controls retain IDs. All magnitudes and all 7,150 point pair records are included.

S8 Reproducibility and access

Table S17: Versions and access status of the source objects.

Object	Manifest / SHA-256	Access
Clean corpus	data/corpus_clean_v1/embedding_manifest.json d8b48bccda4196063535ed6a8b3029481feddb7642e4a3f9eb146a772c71cde5	500,000 IDs; local
Embedding audit	research/embedding_final_audit_2026-09-17/full_verification.json 6e43e878f4409b73b07a74b2ad8f4d9bdf03022c98e4aad91ba7424ea4c47770	Ten models; local vectors
Core report	reports/analysis_v1/final/catalog.json 1309c1458b441962b9b2f5400c074f67c54a0c108a3e04decf84bcc709687bac	Tables / figures; public baseline
Expanded report	reports/robustness_v2/final/catalog.json 2e4ab4b3e0a9445f8456667979162998529da16498e2a26d3778fb1f19415437	Tables / figures; public baseline
Case atlas	reports/prepaper_v1/catalog.json 0e75af65da14a98cc86daf9a30305a858e2034bc922d3bb208b9bcbb30ebcfe8	Examples; public baseline
Morphology	reports/morphology_pilot_v1/catalog.json 8584ef7d693a1d28bc2f4cab60d47f8ab7c7edd5a2f9a61a44f0e11582ac8e2d	Local extension; no data deposit
Field pairs	reports/field_pair_summary_v1/catalog.json 8c009fcadc76702ca84a22389f863f64e1cd068df6be8b6f719cf3715f62305d	Local extension; no data deposit
Writing blueprint	research/writing_blueprint_2026-09-18/closure_audit.json 6f1890922b57893cbb7dcbff23e1e5ad84dae233f6bd4d61f0fbcfa4b17acfa8	Local editorial plan

Hashes identify the listed manifest/catalogue bytes, not the whole folder directly; each catalogue records its own assets. Local reproducibility is broader than public access. No persistent data deposit, general licence or new public release is claimed. Corpus, vectors, model weights, API keys and third-party full texts are excluded from the manuscript bundle. The two PDFs can be rebuilt from the included LaTeX and figure/table assets without the corpus. Full experimental reproduction still requires the separately authorised data-release plan.

S8.1 Rebuilding this display package

The repository command `./manuscript/build.sh` compiles the included LaTeX into the main document and this supplement. It does not run models, searches or simulations. The compiler may download standard TeX support files on its first use. Re-exporting plots and table summaries is a separate presentation command, documented in the package README.

S8.2 Access limits of this draft

The package provides the editable manuscript, figures, descriptive tables and their source summaries. It excludes corpus text, vectors, model weights, secrets and copies of third-party full texts. It can therefore reproduce the two displayed PDFs without reproducing the entire scientific pipeline. The latter requires the frozen inputs and execution materials listed above. A versioned public release and any necessary data permissions/licence must be settled before claiming full public reproducibility. No persistent deposit has yet been made.

The author has confirmed sole authorship as an independent researcher, no external funding and no competing interests. The current prose and scientific interpretations remain a draft for author review. Correspondence details and final journal formatting should be completed at submission; they have not been invented here.

References

- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 3519–3529. PMLR, 09–15 Jun 2019. URL <https://proceedings.mlr.press/v97/kornblith19a.html>.
- Alex Murphy, Joel Zylberberg, and Alona Fyshe. Correcting biased centered kernel alignment measures in biological and artificial neural networks. arXiv:2405.01012, 2024. URL <https://arxiv.org/abs/2405.01012>. Preprint.
- Olivier Roy and Martin Vetterli. The effective rank: a measure of effective dimensionality. In *15th European Signal Processing Conference (EUSIPCO)*, pages 606–610. EURASIP, 2007. URL <https://www.eurasip.org/Proceedings/Eusipco/Eusipco2007/Papers/a5p-h05.pdf>. Open copy also archived at doi:10.5281/zenodo.40328.
- William Rudman, Nate Gillman, Taylor Rayne, and Carsten Eickhoff. IsoScore: Measuring the uniformity of embedding space utilization. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Findings of the Association for Computational Linguistics: ACL 2022*, pages 3325–3339, Dublin, Ireland, may 2022. Association for Computational Linguistics. doi:10.18653/v1/2022.findings-acl.262. URL <https://aclanthology.org/2022.findings-acl.262/>.

Ulrike von Luxburg. A tutorial on spectral clustering. *Statistics and Computing*, 17(4):395–416, 2007. ISSN 0960-3174. doi: 10.1007/s11222-007-9033-z.

Alex H Williams. Equivalence between representational similarity analysis, centered kernel alignment, and canonical correlations analysis. In Marco Fumero, Clementine Domine, Zorah Lähner, Donato Crisostomi, Luca Moschella, and Kimberly Stachenfeld, editors, *Proceedings of UniReps: the Second Edition of the Workshop on Unifying Representations in Neural Models*, volume 285 of *Proceedings of Machine Learning Research*, pages 10–23. PMLR, 14 Dec 2024. URL <https://proceedings.mlr.press/v285/williams24a.html>.

Alex H Williams, Erin Kunz, Simon Kornblith, and Scott Linderman. Generalized shape metrics on neural representations. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 4738–4750. Curran Associates, Inc., 2021. URL https://proceedings.neurips.cc/paper_files/paper/2021/file/252a3dbaeb32e7690242ad3b556e626b-Paper.pdf.